



# Kubernetes Deep Dive

---


## Course Overview

```
mirror_mod = modifier_ob.  
# Set mirror object to mirror  
mirror_mod.mirror_object =  
# operation == "MIRROR_X":  
mirror_mod.use_x = True  
mirror_mod.use_y = False  
mirror_mod.use_z = False  
# operation == "MIRROR_Y":  
mirror_mod.use_x = False  
mirror_mod.use_y = True  
mirror_mod.use_z = False  
# operation == "MIRROR_Z":  
mirror_mod.use_x = False  
mirror_mod.use_y = False  
mirror_mod.use_z = True  
  
# selection at the end -add  
obj.select= 1  
obj.select=1  
context.scene.objects.active  
("Selected" + str(modifier_ob.  
mirror_ob.select = 0  
= bpy.context.selected_object  
data.objects[one.name].select  
  
print("please select exactly  
  
-- OPERATOR CLASSES ----  
  
types.Operator):  
X mirror to the selected  
object.mirror_mirror_x"  
mirror X"  
  
context):  
context.active_object is not
```

# Deep Dive

- Builds upon the Fundamentals course
- Requires knowledge of Kubernetes architecture and the basic building blocks of **Pods**, **Deployments**, **ReplicaSets**, **ConfigMaps**, **Secrets**, ...

Deep dive ≠ under the hood  
look at Kubernetes



## Geert Baeke

Cloud-native architect and Microsoft Azure  
Incubator and Cloud-Native Architect @  
St-Niklaas, Flemish Region, Belgium · [Contact](#)



Home



Explore



Notifications



Messages



Bookmarks



Lists



Profile



## Geert Baeke

7,926 Tweets



Cloud-native architect  
Blogger - Speaker



Edit profile

## Geert Baeke

@GeertBaeke

| Technology at the intersection of cloud-native infrastructure and applications

| Microsoft Azure MVP

| Check out my YouTube channel: [youtube.com/geertbaeke](https://youtube.com/geertbaeke)

Belgium [blog.baeke.info](https://blog.baeke.info) Joined March 2008

# What's in this course?

---

## AKS Networking

- Ingress and egress control
- Service Mesh

## Advanced pod concepts

- Init containers
- Multi-container pods
- Sidecars
- Dapr

## Advanced release concepts

- Kustomize
- Helm
- Progressive Delivery
- GitOps

## Security, authentication and authorization

- RBAC
- Secrets handling
- Pod identity





# AKS Networking

---

Kubernetes Deep Dive

```
mirror_mod = modifier_ob.  
#set mirror object to mirror  
mirror_mod.mirror_object =  
operation == "MIRROR_X":  
    mirror_mod.use_x = True  
    mirror_mod.use_y = False  
    mirror_mod.use_z = False  
operation == "MIRROR_Y":  
    mirror_mod.use_x = False  
    mirror_mod.use_y = True  
    mirror_mod.use_z = False  
operation == "MIRROR_Z":  
    mirror_mod.use_x = False  
    mirror_mod.use_y = False  
    mirror_mod.use_z = True  
  
#selection at the end -add  
mirror_ob.select= 1  
modifier_ob.select=1  
context.scene.objects.active  
("Selected" + str(modifier_ob.  
mirror_ob.select = 0  
= bpy.context.selected_object  
data.objects[one.name].select  
print("please select exactly  
-- OPERATOR CLASSES ----  
  
types.Operator):  
X mirror to the selected  
object.mirror_mirror_x"  
mirror X"  
  
context):  
context.active_object is not
```

# What's in this module?

---

Azure Virtual  
Networks

AKS and  
VNET  
Integration

Egress  
control

Network  
Policies

Service  
Mesh



# Azure Virtual Networks

---

# Azure Virtual Network Concepts

---



Virtual  
Network



Subnet



Load balancer



Routing table



Network  
Security Group

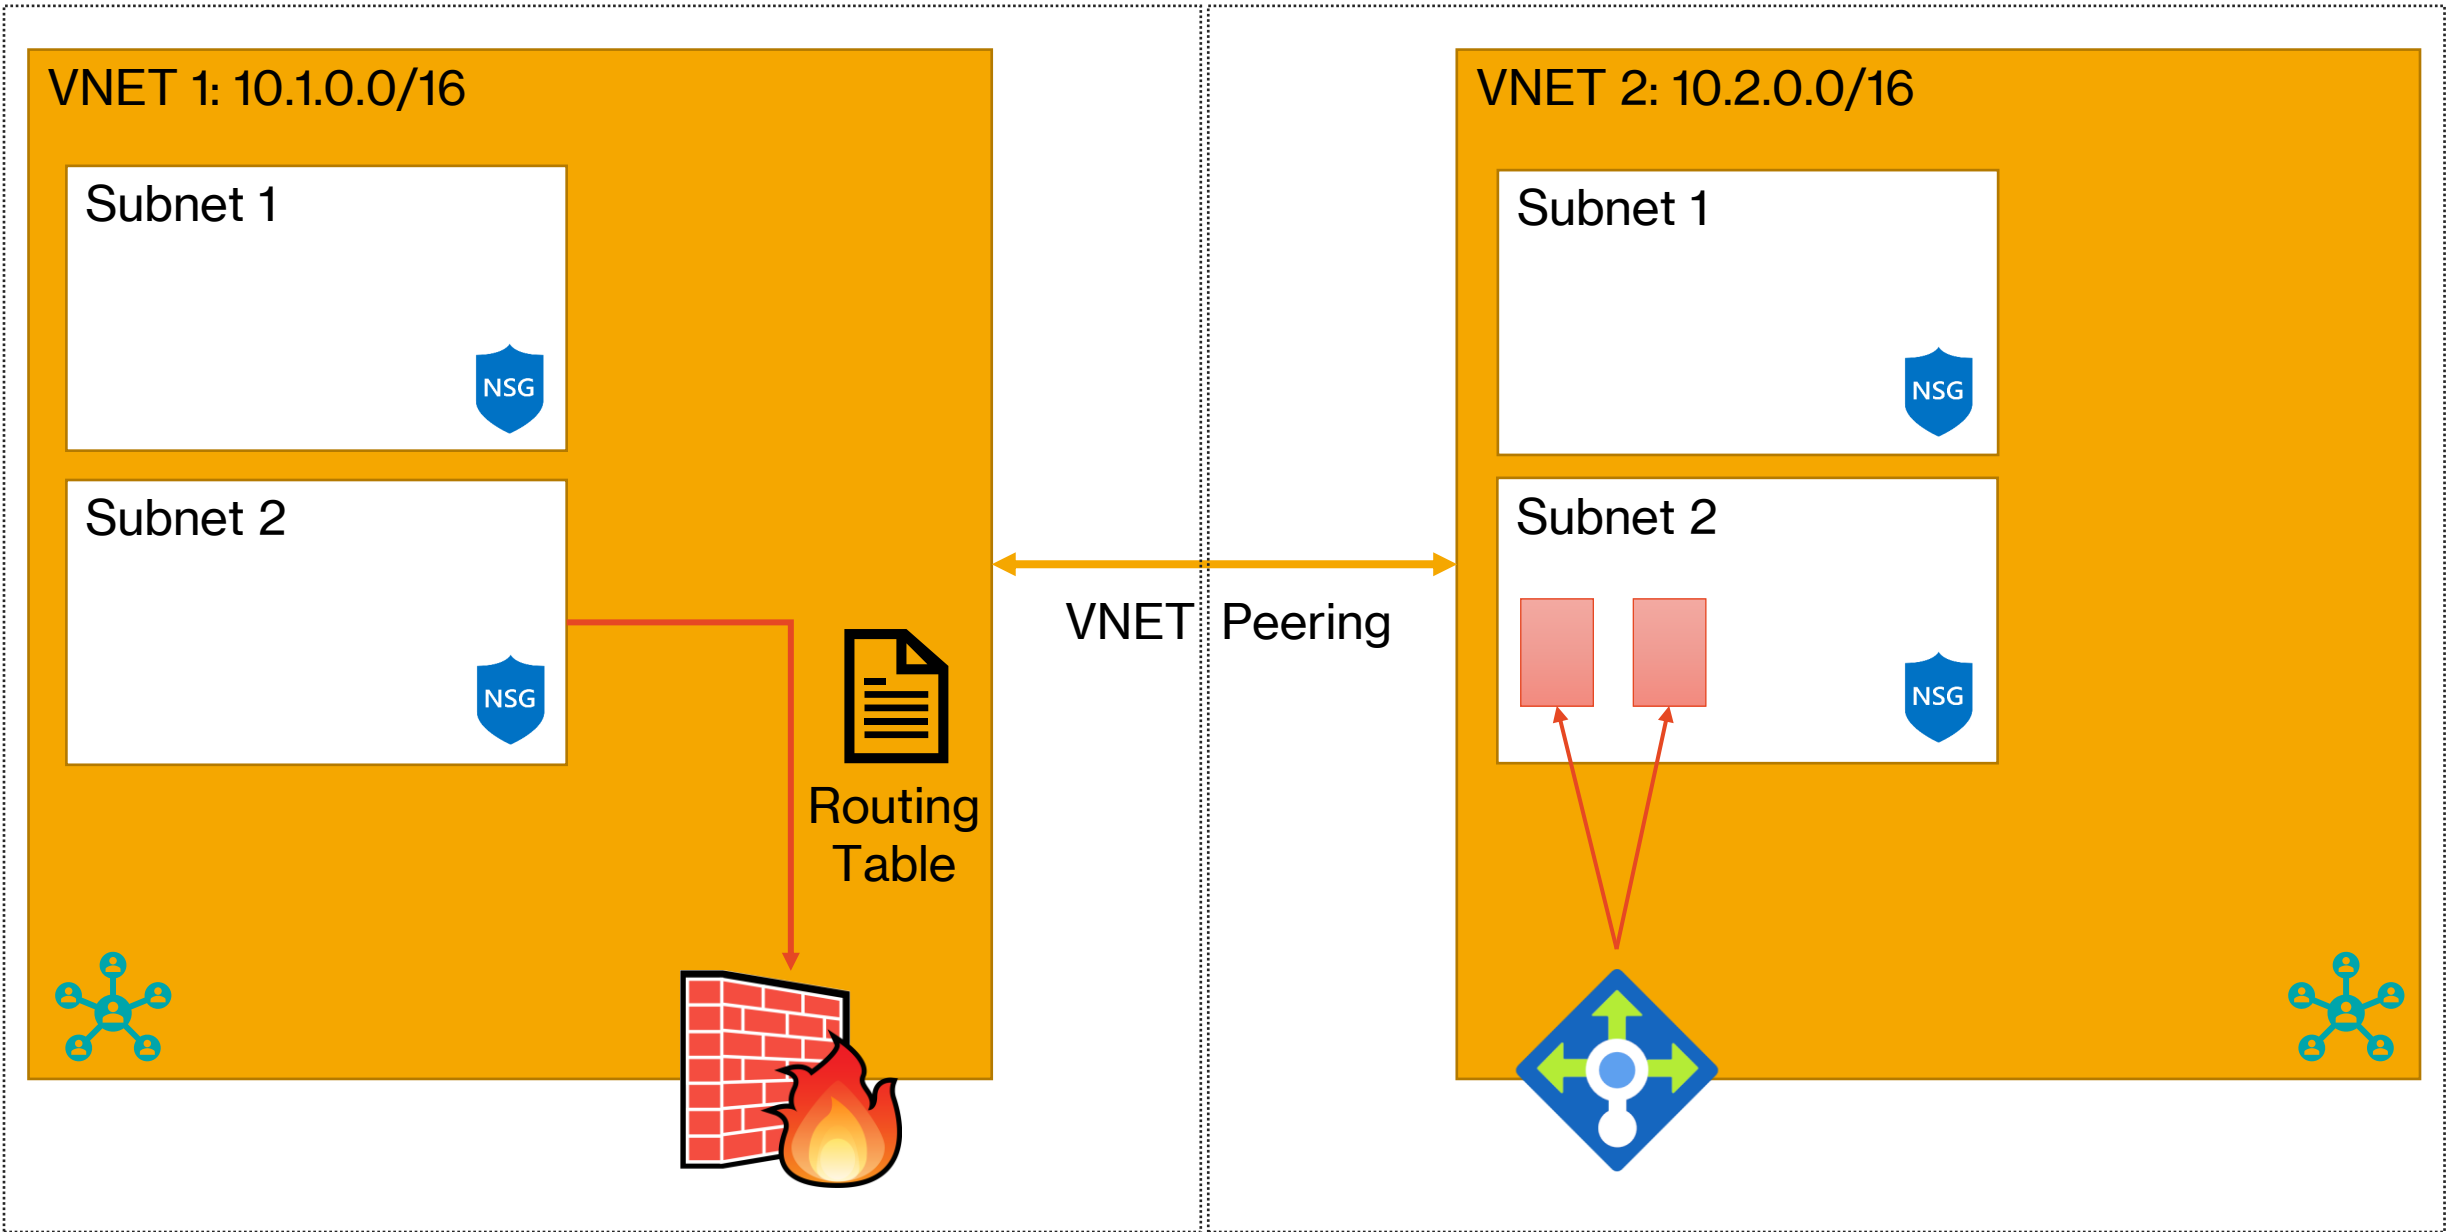


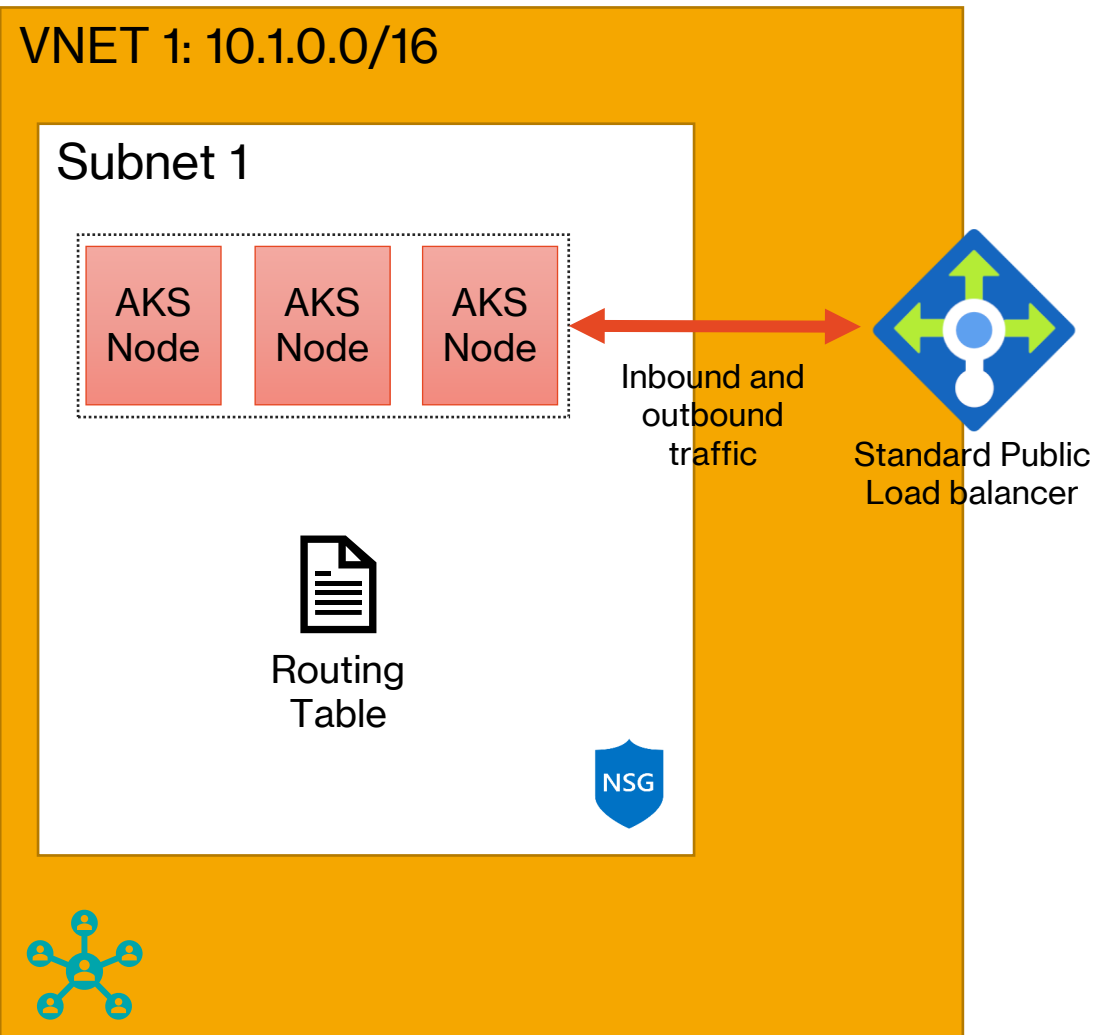
Azure Firewall



Subscription A

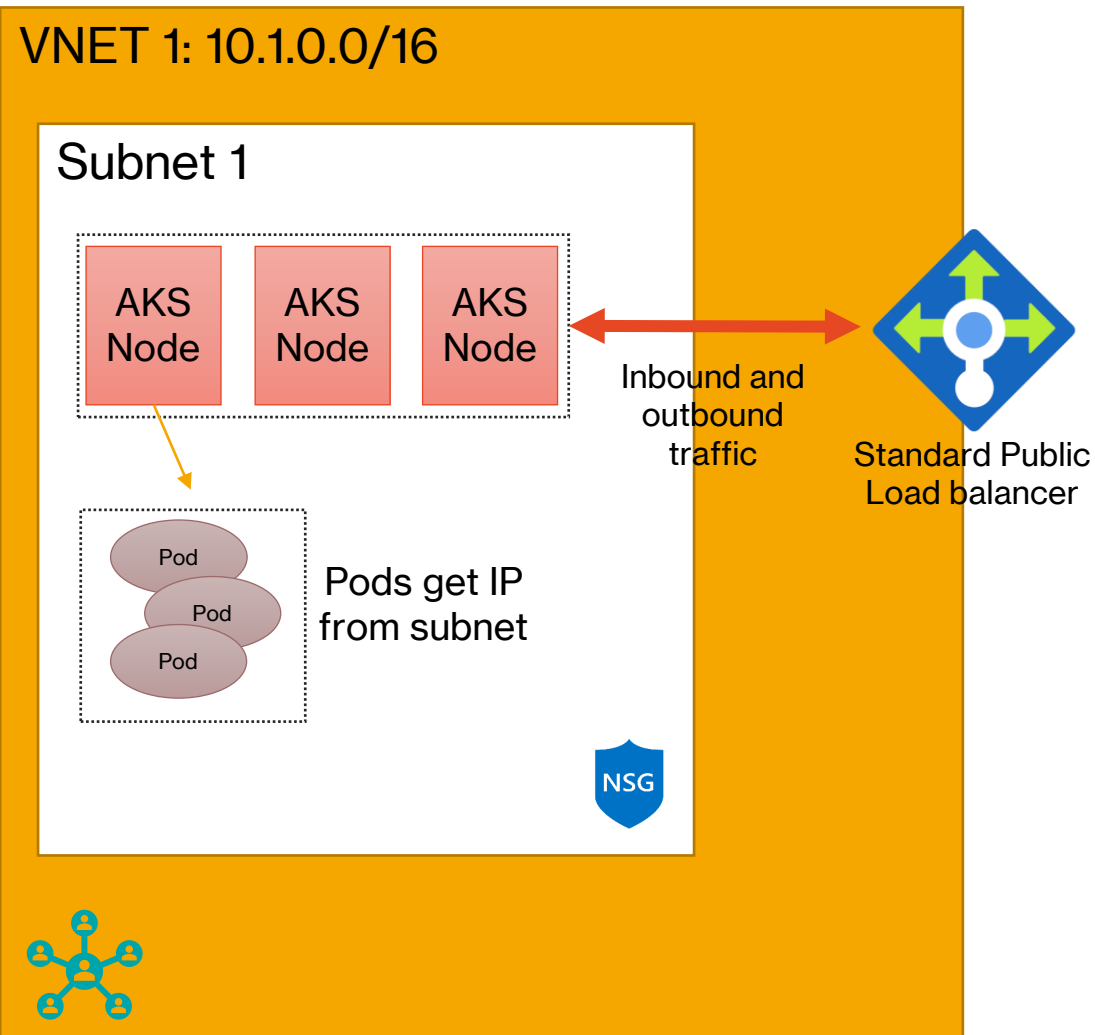
Subscription B





## AKS Kubelet networking

- Node-only routable IP address
- Pods use NAT
- 110 pods per node
- 400 nodes per cluster (400 routes in a UDR)
- Supports Calico Network Policies
- Does not support:
  - Virtual nodes with ACI
  - Azure Network Policies



## AKS Azure CNI networking


- Every pod gets a routable IP address
- IP blocks reserved per node
- 30 default, up to 250 pods per node
- 1000 nodes per cluster (100 per node pool)
- Supports:
  - Virtual Nodes
  - Calico and Azure Network Policies

# Networking preview features


- Dynamic allocation of IPs and enhanced subnet support
- IPs are dynamically allocated from a pod subnet
- Each node pool can use a different subnet for nodes and pods
- Scale node and pod subnet independently
- Works with network policies (Azure and Calico)

VNET 1: 10.1.0.0/16

Subnet 1



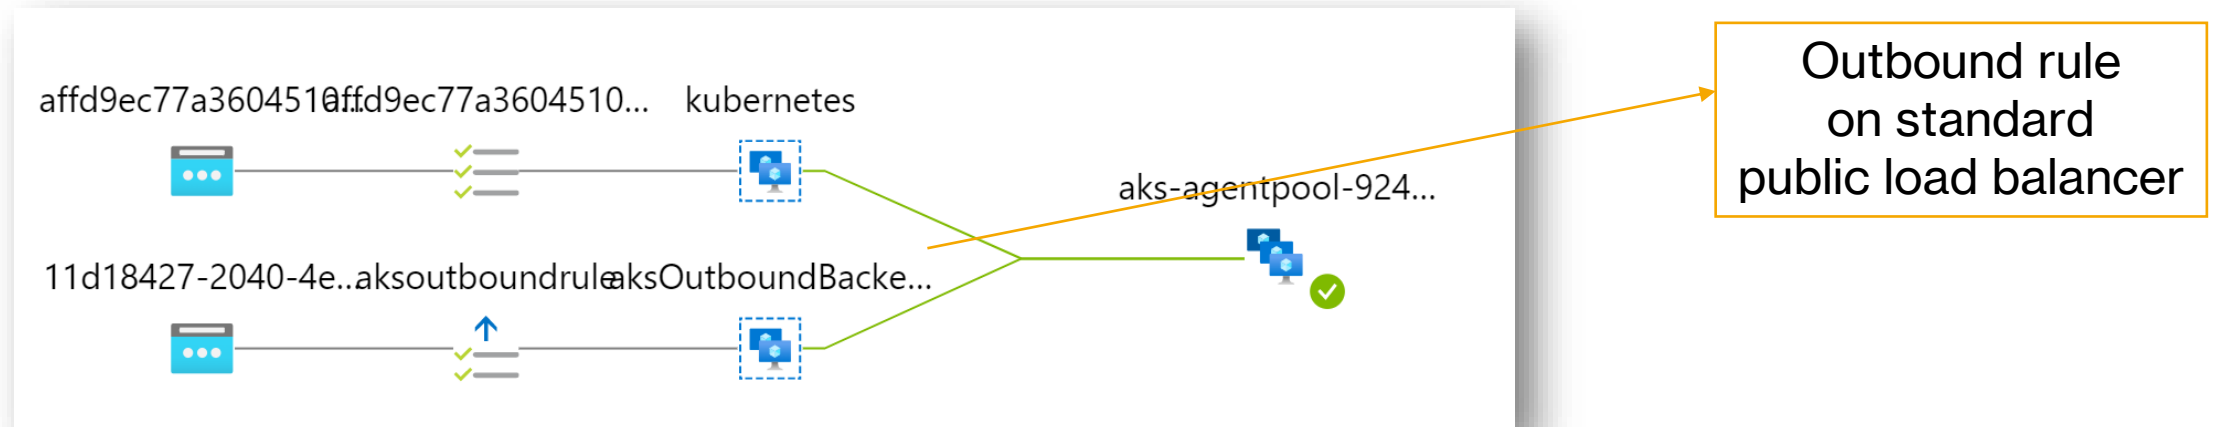
Subnet 2



Pods get IP  
from subnet 2

# Controlling outbound traffic

- AKS clusters have unrestricted outbound Internet access
- Allows nodes and pods access to external resources as needed
- AKS needs access to several external resources for proper operation (via FQDNs without static IPs)







# Preview: Managed NAT Gateway

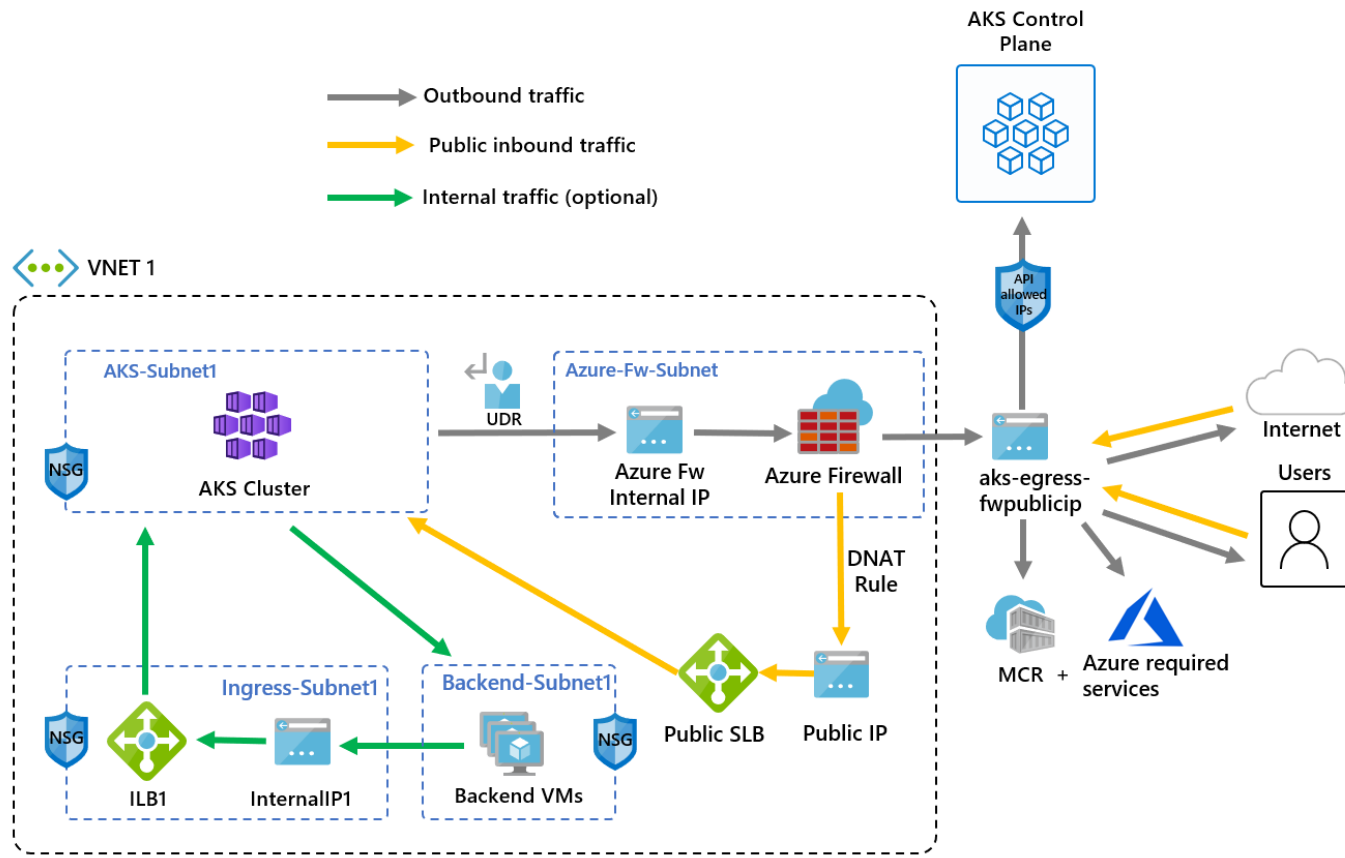
---

- Support up to 64000 UDP and TCP traffic flows per IP address
  - maximum 16 IP addresses
- New **--outbound-type:** managedNATGateway
  - set # of IPs with --nat-gateway-managed-outbound-ip-count

# Preview: HTTP Proxy support

- Configure AKS nodes to use a proxy to access the Internet
- Requires a proxy configuration file (JSON) and the use of the **--http-proxy-config** parameter
- Several limitations at the moment:
  - Linux only
  - Does not support monitoring addon
  - User/password authentication not supported

# Using Azure Firewall

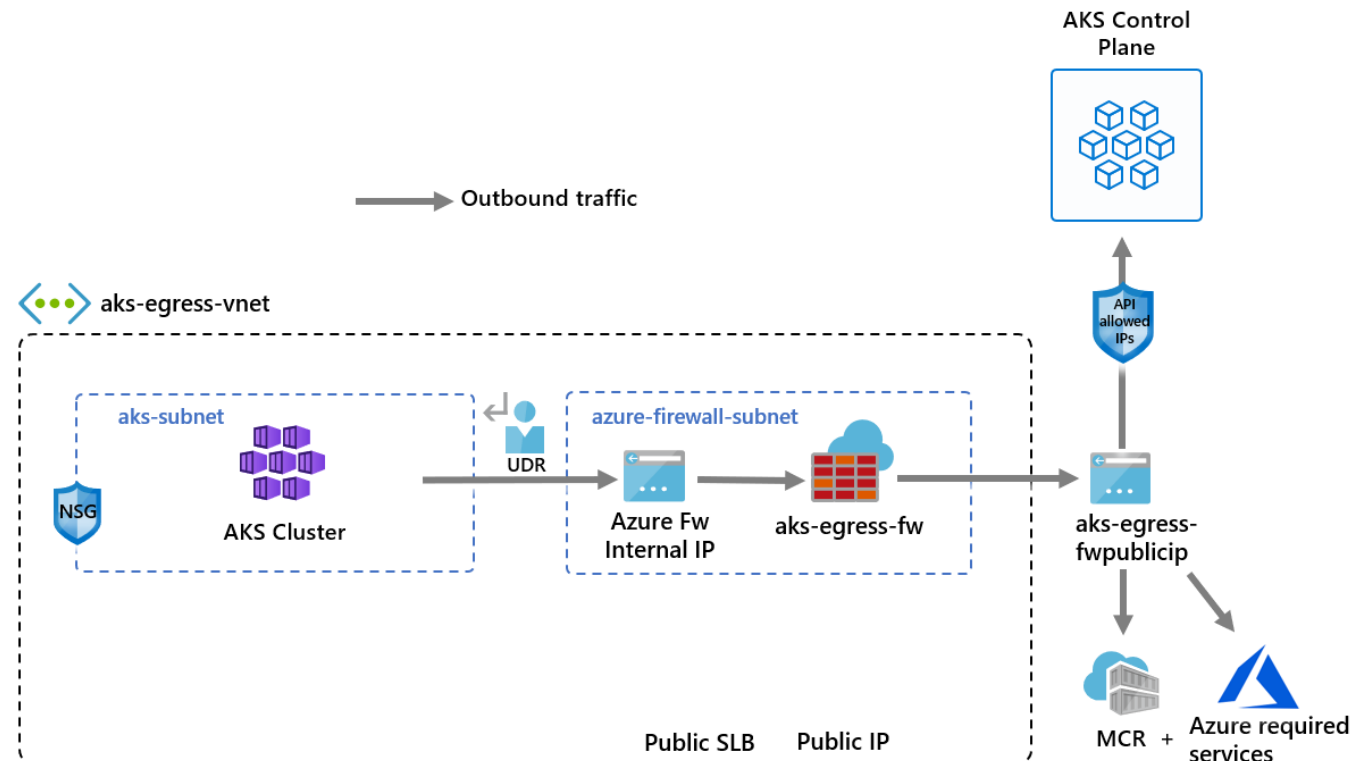


⚠ Azure Firewall has an *AzureKubernetesService* FQDN tag to simplify the configuration of outbound rules

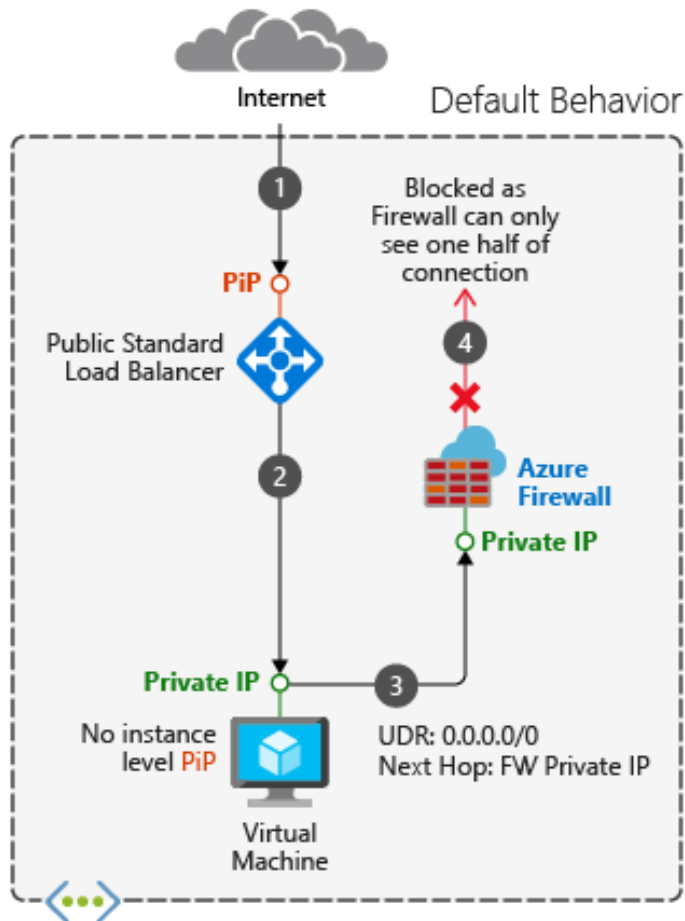
# Setting the outboundType

AKS clusters have an **outboundType** setting:

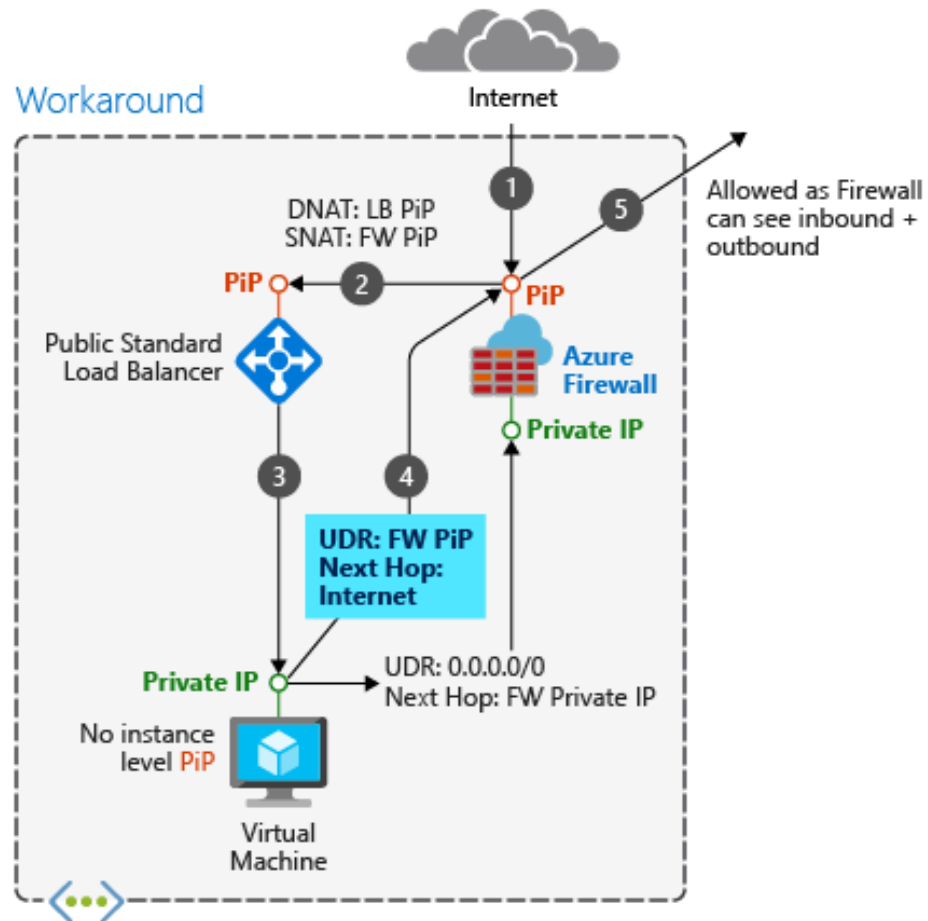
- loadBalancer
- userDefinedRouting
- managedNATGateway (preview)



# Beware of asymmetric routing



## Workaround












# Cluster Deployment Tips

---

## Create Kubernetes cluster




Choose your scenario to view and apply the recommended configurations suited to your needs. The settings in the table below will be updated to the specified values based on your selection. All other cluster settings will remain unchanged. [Learn more](#)

	 <b>Standard (\$\$)</b> <ul style="list-style-type: none"> <li>Best if you're not sure what to choose.</li> <li>Works well with most applications.</li> </ul> <input checked="" type="radio"/>	 <b>Dev/Test (\$)</b> <ul style="list-style-type: none"> <li>Best for experimenting with AKS or deploying a test app.</li> </ul> <input type="radio"/>	 <b>Cost-optimized (\$)</b> <ul style="list-style-type: none"> <li>Best for reducing costs on production workloads that can tolerate interruptions.</li> </ul> <input type="radio"/>	 <b>Batch processing (\$\$\$)</b> <ul style="list-style-type: none"> <li>Best for machine learning, compute-intensive, and graphics-intensive workloads.</li> <li>Suited for apps requiring fast scale-up and scale-out.</li> </ul> <input type="radio"/>	 <b>Hardened access (\$\$\$)</b> <ul style="list-style-type: none"> <li>Best for large enterprises that need full control of security and stability.</li> </ul> <input type="radio"/>
System node pool node size ⓘ	DS2_v2 ⓘ	B4ms ⓘ	B4ms ⓘ	D4s_v3 ⓘ	D4s_v3 ⓘ
User node pool node size ⓘ	-	-	B4ms ⓘ	NC6s_v2 ⓘ	D4s_v3 ⓘ
Cluster autoscaling ⓘ	✓	-	✓	✓	✓
Private cluster ⓘ	-	-	-	-	✓
Availability zones ⓘ	✓	-	-	-	✓
Azure Policy ⓘ	-	-	-	-	✓
Azure Monitor ⓘ	✓	-	-	✓	✓

# AKS Deploy Helper

See <https://azure.github.io/Aks-Construction>

Use for learning; not recommended to use as-is in production


**AKS Deploy helper**

Provide the requirements of your AKS deployment to generate the assets to create a full operational environment, incorporating best-practices guidance. For documentation, and CI/CD samples - please refer to our [GitHub Repository](#)

Enterprise Scale ⓘ  
☐ No


### Operations Principles

☐ Simplest bare-bones cluster



Just Kubernetes please, I will make decisions later


☐ I prefer control & community open source solutions



Use proven, open source projects for my Kubernetes operational environment, and self-manage my clusters upgrades and scaling

- Manual Upgrades
- Manual Scaling
- Contour Ingress ([docs](#))
- Prometheus/Grafana Monitoring ([docs](#))
- DockerHub container registry

☒ I want a managed environment




I'd like my cluster to be auto-managed by Azure for upgrades and scaling, and use Azure provided managed addons to create an full environment with the minimum of operational requirements

- Cluster auto-scaler ([docs](#))
- Cluster auto-upgrades ([docs](#))
- Azure Monitor for Containers ([docs](#))
- Azure Container Registry
- Azure AppGateway Ingress ([docs](#))


### Security Principles

☐ Simple cluster with no additional access limitations



Simplest option for experimenting with kubernetes, or clusters with no sensitive data


☒ Cluster with additional security controls



Good option for implementing recommended minimum security controls for regular environments

- AAD Integration ([docs](#))
- AUDIT Pod security baseline standards ([docs](#))
- East-West traffic control ([docs](#))
- Authorized IP address ranges ([docs](#))
- Restrict dependencies with Service Endpoints \*\* (ACR preview) ([docs](#))

☐ Private cluster with isolating networking controls



Best option for high-secure, regulated environments or sensitive data requirements.

ⓘ **WARNING:** most complex environment option to operate

- AAD Integration ([docs](#))
- ENFORCE Pod security baseline standards ([docs](#))
- East-West traffic control ([docs](#))
- Private Cluster ([docs](#))
- Restrict dependencies with Private Link ([docs](#))
- Restrict egress with Azure Firewall ([docs](#))
- Store Kubernetes Secrets in Azure KeyVault ([docs](#))

# Uptime SLA & HA

- Use the Uptime SLA in production and availability zones
- 99.95% vs. 99.5% SLO

```
gbaeke ~ $ kubectl describe nodes | grep -e "Name:" -e "failure-domain.beta.kubernetes.io/zone"
Name:          aks-agentpool-13192139-vmss000003
failure-domain.beta.kubernetes.io/zone=westeurope-1
Name:          aks-agentpool-13192139-vmss000004
failure-domain.beta.kubernetes.io/zone=westeurope-2
```

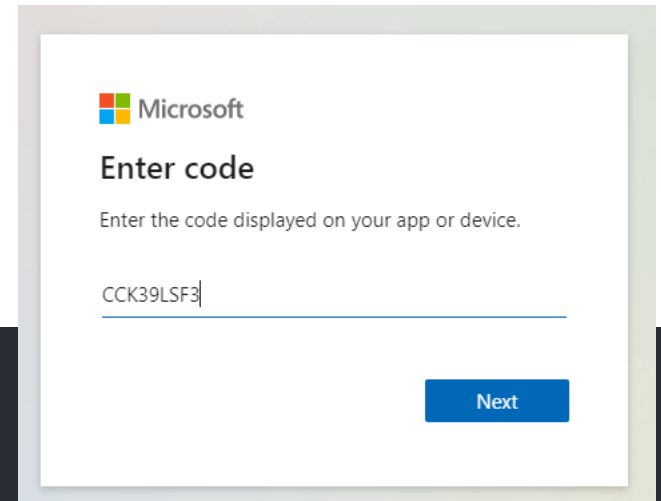
label on each node  
use **kubectl describe node <nodename>** to see all labels

⚠ k8s 1.17.0 and later: use of **topology.kubernetes.io/zone** label  
(other is deprecated)

# Azure Active Directory Authentication

- Use AKS-managed option
- Add one or more Azure AD groups

```
gbaeke ~ $ az aks get-credentials -n myCluster2 -g rg-aks
The behavior of this command has been altered by the following extension:
Merged "myCluster2" as current context in /home/gbaeke/.kube/config
gbaeke ~ $ kubectl get nodes
To sign in, use a web browser to open the page https://microsoft.com/devicelogin and enter the code CCK39L
SF3 to authenticate.
```



Microsoft

Enter code

Enter the code displayed on your app or device.


CCK39LSF3

Next



# Cluster infrastructure authentication

- AKS comes with a cloud provider that needs to deploy additional resources:
  - Load balancers
  - Managed disks
- Cloud provider needs an identity
  - Service principal – or –
  - Managed Identity
- Use managed identity 👍
  - Wrapper around service principals
  - No need to update secrets




# Control plane and kubelet identities

---



## Control plane

Create load balancers  
Create AKS managed public Ips  
Cluster autoscaler  
Azure Disk and Azure File



## kubelet

Authentication with Azure Container  
Registry (ACRPull role)

# MC\_rg-aks\_clu-vn\_westeurope | Access control (IAM) ...

Search (Ctrl+ /)

- Overview
- Activity log
- Access control (IAM)
- Tags
- Resource visualizer
- Events

## Settings

- Resource costs
- Deployments
- Security
- Policies
- Properties
- Locks

## Monitoring

- Insights (preview)

+ Add   ↓ Download role assignments   ≡ Edit columns   ↻ Refresh   ✕ Remove   👤 Got feedback?

Check access   Role assignments   Roles   Deny assignments   Classic administrators

### Number of role assignments for this subscription ⓘ

56

2000

Search by name or email






Type : All

Role : All

Scope : All scopes

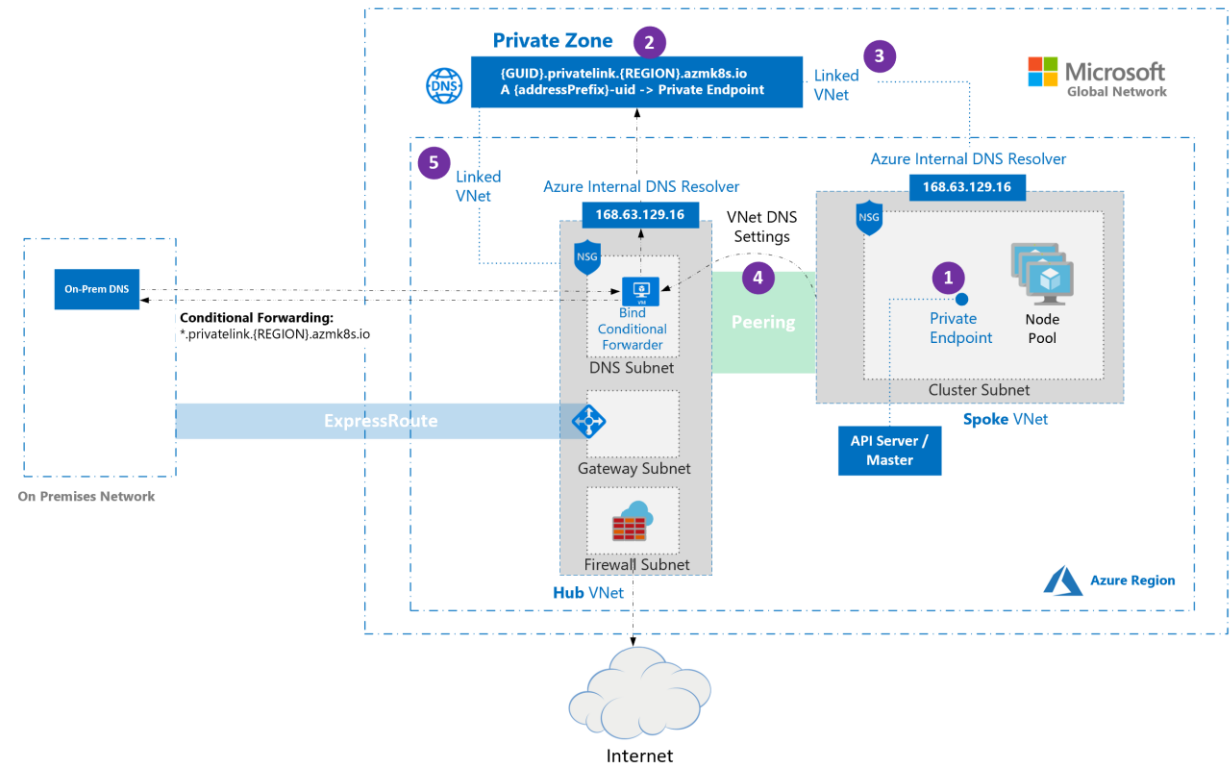
Group by : Role

12 items (2 Users, 1 Groups, 5 Service Principals, 2 Unknown, 2 Managed Identities)

<input type="checkbox"/> Name	Type	Role	Scope	Condition
▼ Contributor				
<input type="checkbox"/>  <a href="#">msi-connection-aks-clu-vn</a> <small>(subscriptions/af1d1b1d-1b1d-1b1d-1b1d-1b1d1b1d1b1d)</small>	User-assigned Managed Identity	<a href="#">Contributor</a> ⓘ	This resource	None
<input type="checkbox"/>  <a href="#">aks-cl-2021-09-21-11-09-09</a>	App	<a href="#">Contributor</a> ⓘ	<a href="#">Subscription</a> (Inherited)	None
<input type="checkbox"/>  <a href="#">clu-vn</a>	App	<a href="#">Contributor</a> ⓘ	This resource	None
<input type="checkbox"/>  <a href="#">identitynot found (2)</a> <small>Users do not find identity.</small>	Unknown	<a href="#">Contributor</a> ⓘ	<a href="#">Subscription</a> (Inherited)	None
<input type="checkbox"/>  <a href="#">subscription-automation</a>	App	<a href="#">Contributor</a> ⓘ	<a href="#">Subscription</a> (Inherited)	None

# Private Cluster

- Gives the API server an IP address in the AKS subnet
- Microsoft exposes API server via **Private Link**, connected to **Private Endpoint** in your own network





DNS

19a824aa-0796-4738-b3a1-06b7e16c5fa1.privatelink.westeurope.azmk8s.io

Private DNS zone

Search (Ctrl+ /)

<<

+ Record set

→ Move

🗑️ Delete zone

🔄 Refresh

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Virtual network links

Properties

Locks

Monitoring

Alerts

Metrics

Automation

Essentials

Resource group (Move) : mc\_rg-aks-clu-pri\_westeurope

Subscription (Move) : Microsoft Azure Sponsorship

Subscription ID : d1d3dadc-bc2a-4495-b8dd-70443d1c70d1

Tags (Edit) : Click here to add tags

You can search for record sets that have been loaded on this page. If you don't see what you're looking for, you can try scrolling to allow more records to load.

Search record sets

19a824aa-0796-4738-b3a1-06b7e16c5fa1.privatelink.westeurope.azmk8s.io | Virtual network link

Private DNS zone

Search (Ctrl+ /)

<<

+ Add

🔄 Refresh

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Virtual network links

Properties

Locks

Monitoring

Alerts

Metrics

Automation

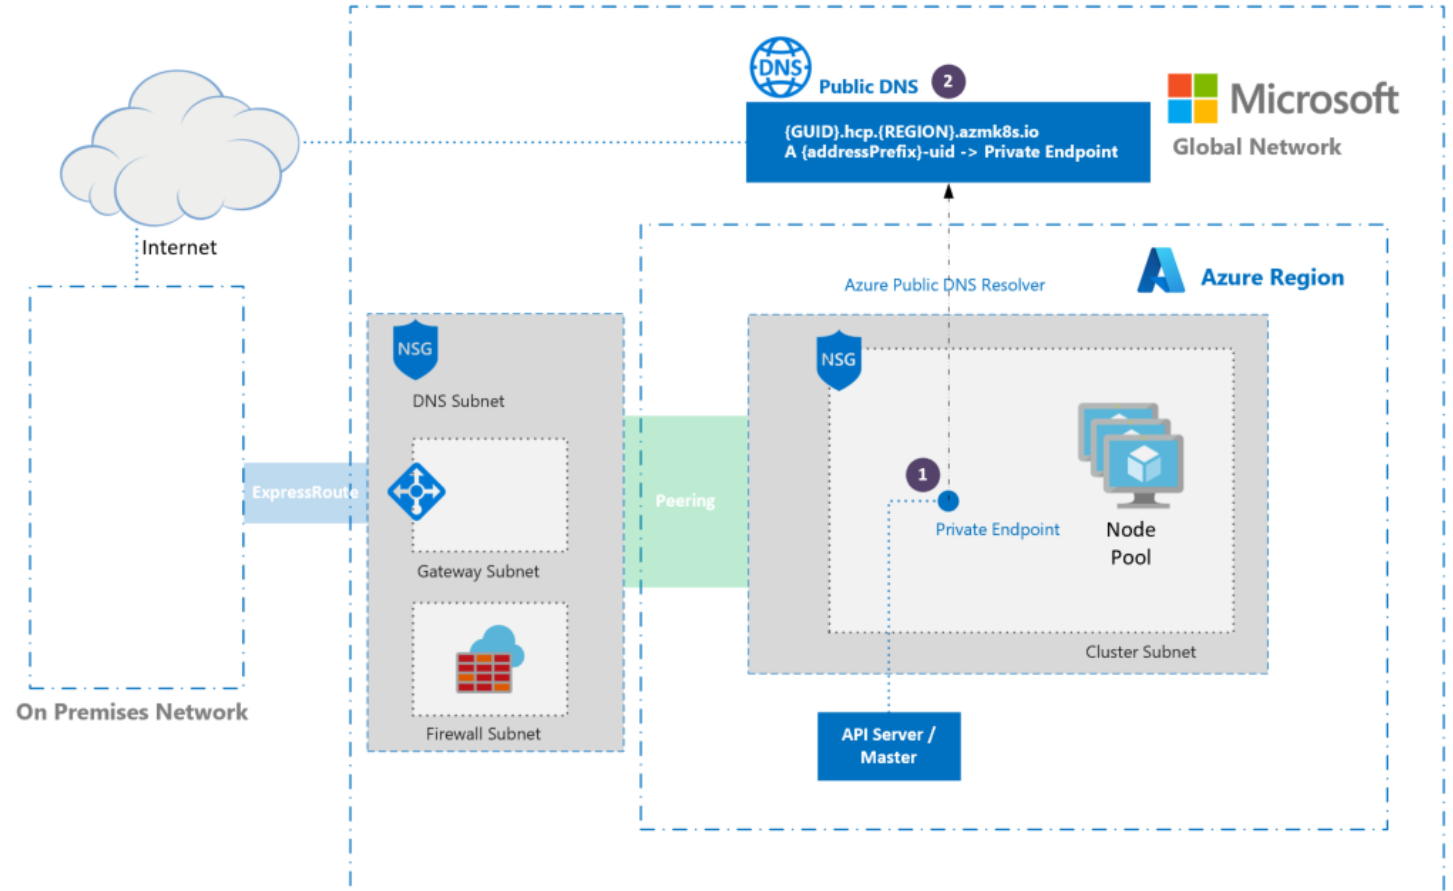
Search virtual network links

Link Name	Link status	Virtual network
clu-pri-dns-843e9087	Completed	rg-aks-vnet



# Public DNS for Private Cluster

- DNS resolution for private clusters can be complex: see <https://docs.microsoft.com/en-us/azure/aks/private-clusters>
- Easiest option: use public FQDN



# Deploy your cluster

- Deploy with AKS Deploy Helper
  - Managed environment
  - Additional security controls
  - Use one node pool (system and user)
  - Turn off Ingress Controller (defaults to AG with WAF)
  - Specify AAD Group ObjectID
    - ⚠ Deploy Helper might fail setting it properly; update Configuration afterwards





# Network Policies

---



# Network Policy

Control  
**ingress** and  
**egress** for  
pods in AKS

Two options:

- Azure Network Policies
- Calico Network Policies

Enable  
network policy  
at cluster  
creation time

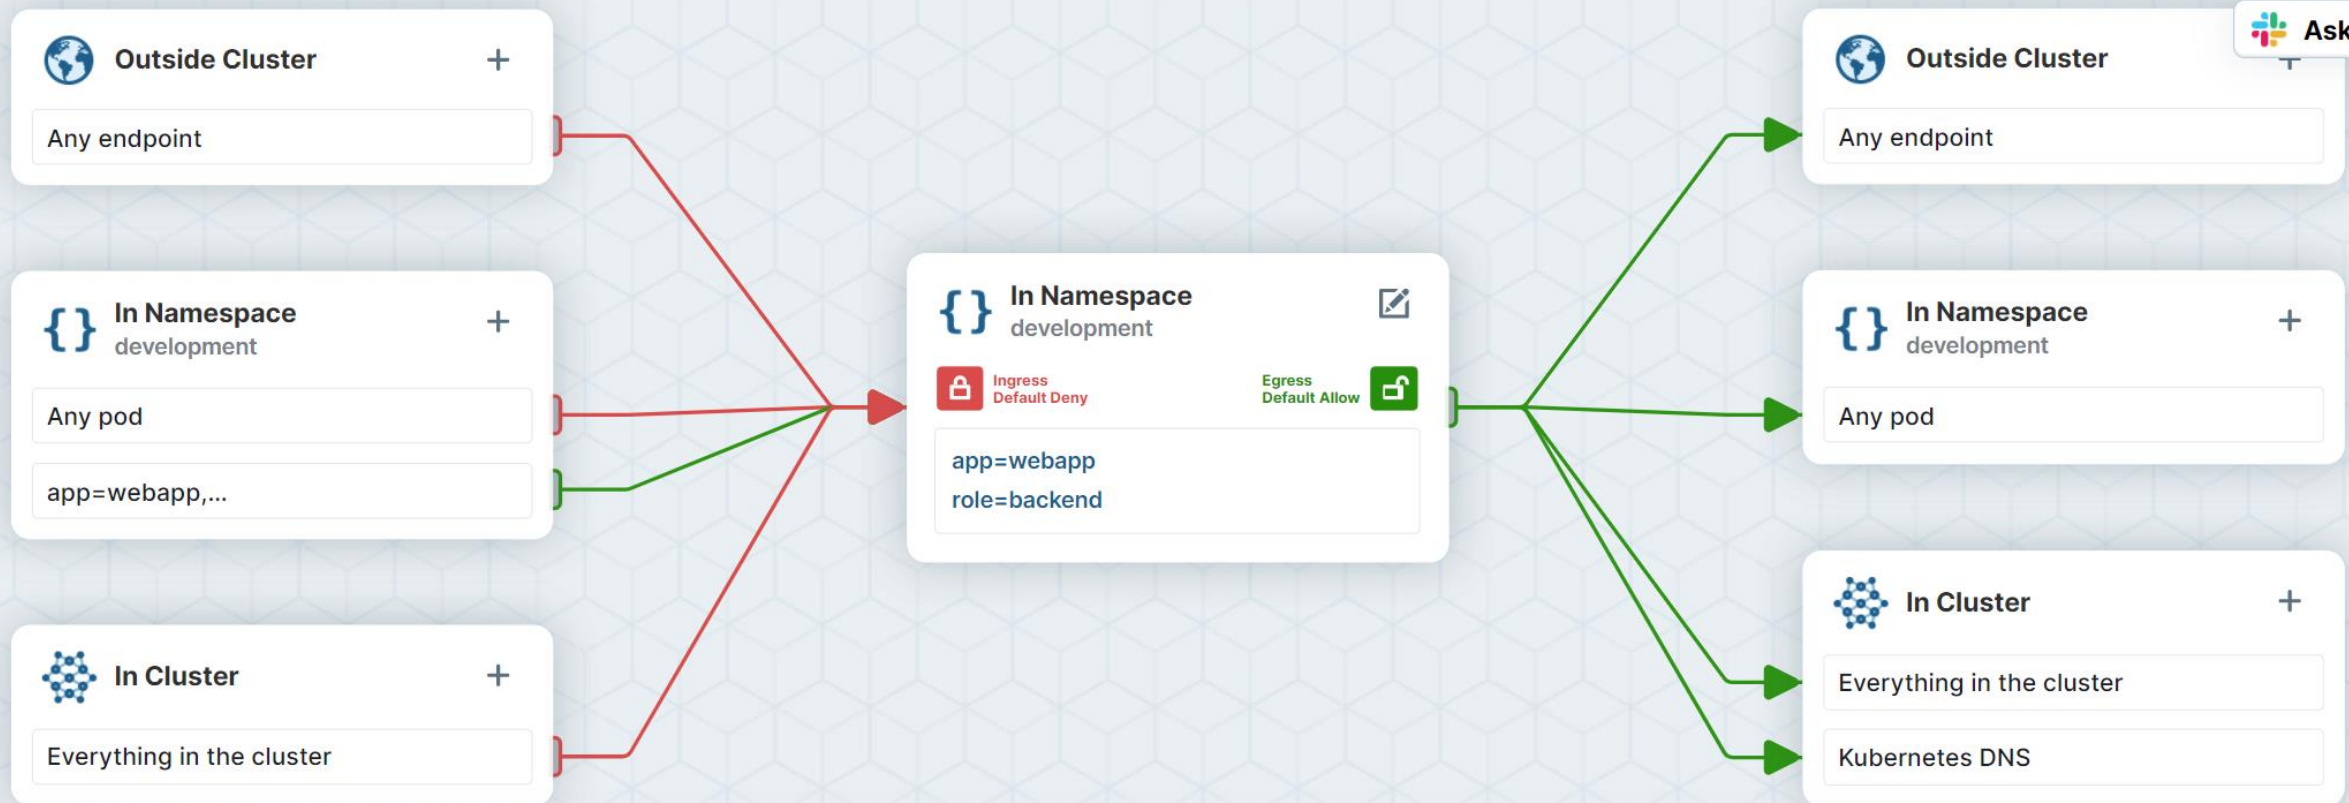
# Example policy

- In the **development** namespace, allow traffic to pods with **backend** labels from pods with **frontend** labels.

⚠ The **frontend** pods can be in any namespace.

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: backend-policy
  namespace: development
spec:
  podSelector:
    matchLabels:
      app: webapp
      role: backend
  ingress:
    - from:
      - namespaceSelector: {}
        podSelector:
          matchLabels:
            app: webapp
            role: frontend
```





Kubernetes Network Policy Cilium Network Policy

Policy Rating ■■■■■

Download

Share

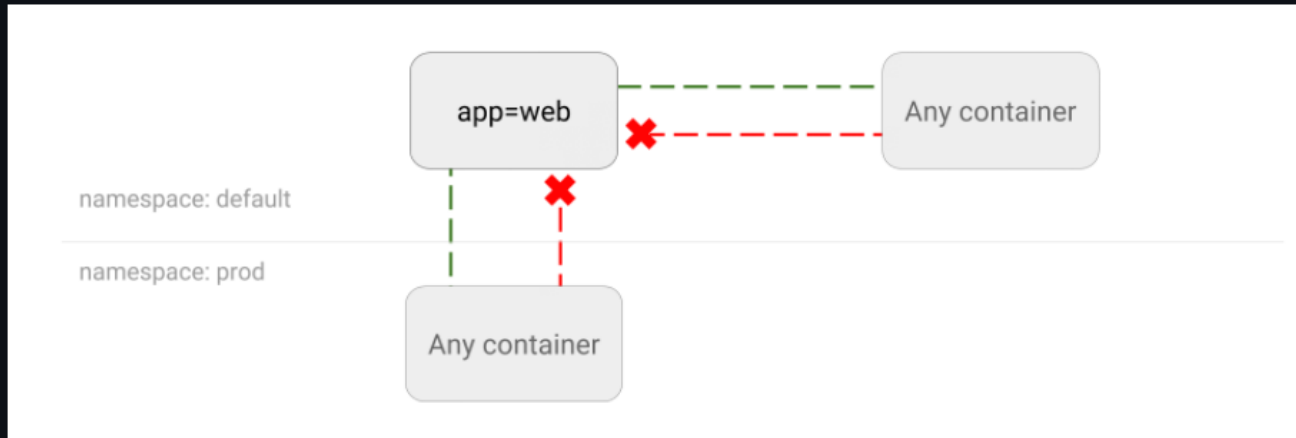
▼ Main tutorial Flows upload

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: backend-policy
5   namespace: development
6 spec:
7   podSelector:
8     matchLabels:
9     app: webapp
```

## Welcome to the Network Policy Editor! <sup>Beta</sup>

This tutorial will teach you how to create a network policy using the Editor. It explains basic network policy concepts and guides you through the steps needed to achieve the desired least- privilege security and zero-trust concepts.

### Step 1. What pods do you want to secure?



*You can get stuff like this with Network Policies...*

## Kubernetes Network Policy Recipes

This repository contains various use cases of Kubernetes [Network Policies](#) and sample YAML files to leverage in your setup. If you ever wondered how to drop/restrict traffic to applications running on Kubernetes, read on.

setup. If you ever wondered how to drop/restrict traffic to applications running on Kubernetes, read on. This repository contains various use cases of Kubernetes [Network Policies](#) and sample YAML files to leverage in your

## Kubernetes Network Policy Recipes



# Calicoctl

- Separate CLI
- Check version of client and server with **calicoctl version**
- Other uses:
  - Implement Calico network policies (namespace)
  - Implement Calico GlobalNetworkPolicy (independent of namespace)
  - Creating NetworkSets and GlobalNetworkSets

```
apiVersion: projectcalico.org/v3
kind: NetworkPolicy
metadata:
  name: calicopol
  namespace: default
spec:
  selector: app == 'superapi'
  ingress:
    - action: Allow
      protocol: TCP
      source:
        selector: app == 'debug'
      destination:
        ports:
          - 8080
```



# Service Mesh

---



# What does a service mesh provide?



## Observability

Collecting metrics  
Golden signals



## Security

Mutual TLS



## Reliability

Retries  
Timeouts  
Load balancing  
Traffic shifting

# Useful in microservice applications

👍 Services use synchronous communication with HTTP or gRPC

👎 Monoliths or services that communicate via queues/topics




Open Service Mesh

---




LINKERD

---



Istio

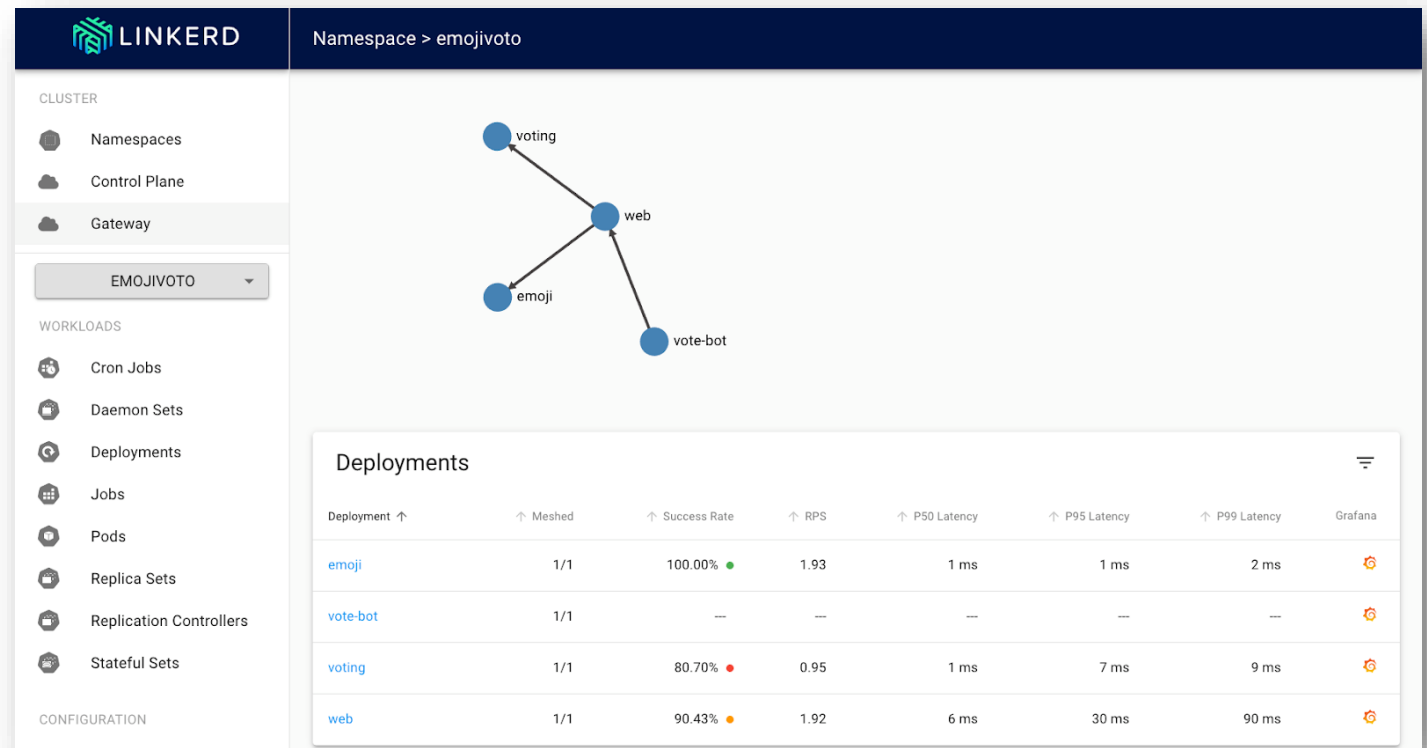


Consul

# Golden Signals

- Latency
- Traffic (RPS)
- Successful responses (SR)


No need to instrument  
your application to obtain  
these signals!



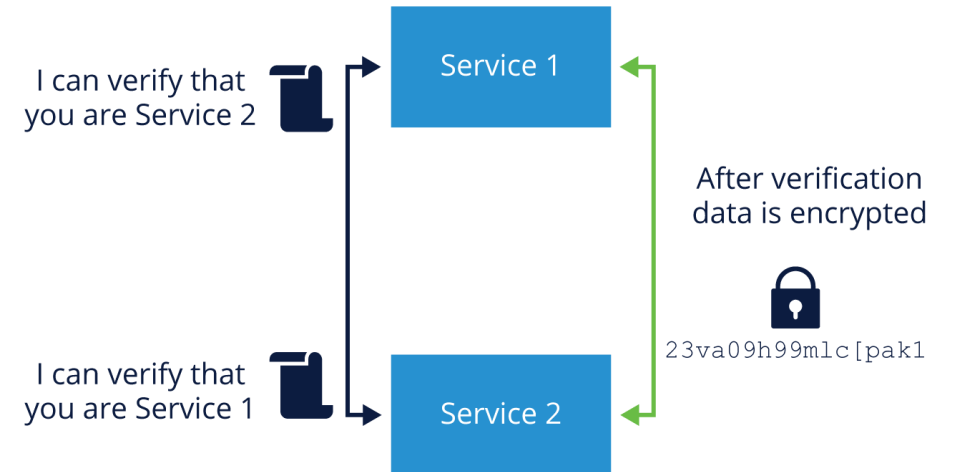
# Security

- Transparent mutual TLS between services
  - ➔ no configuration required
- Provides both **confidentiality** and **authentication**

Without mutual TLS  
data is sent in plaintext



With mutual TLS, identity is  
verified and data is encrypted








# Reliability

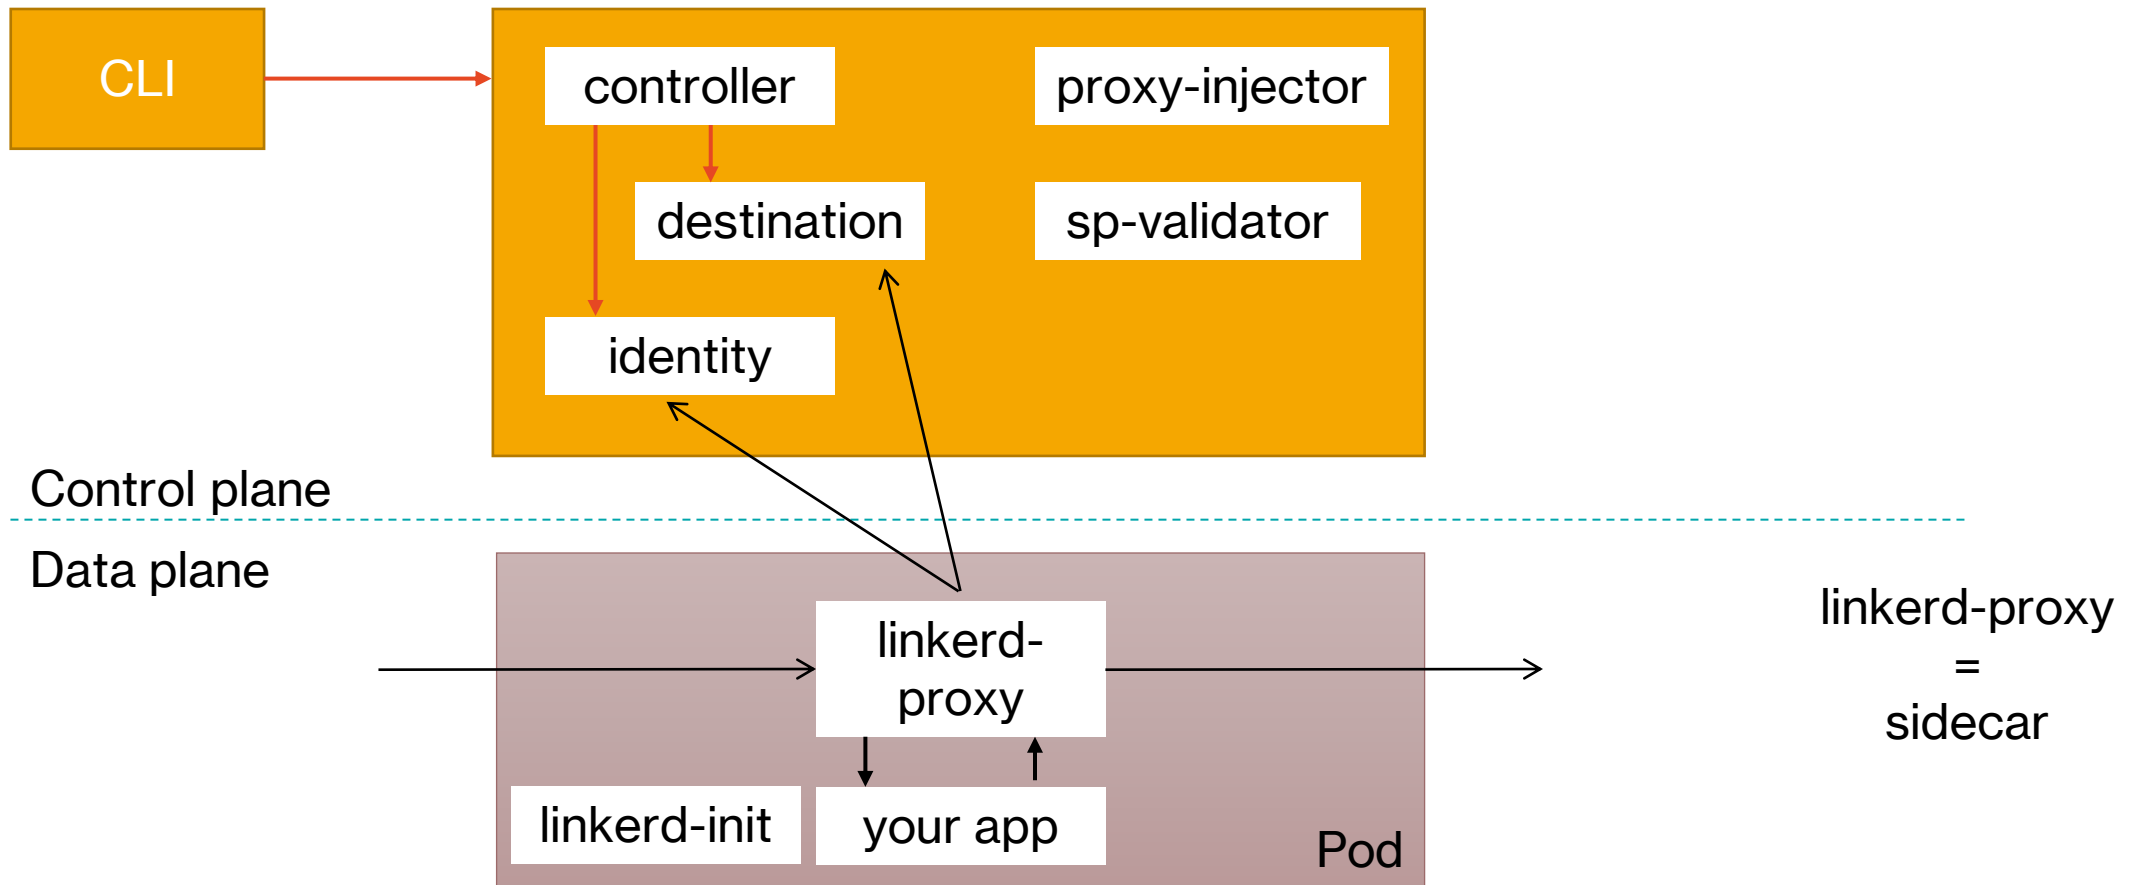
- Retry failed requests
- Configurable timeouts
- Load balancing based on latency
- Traffic shifting

# Service Mesh Interface (SMI)

- Standard set of interfaces that all service mesh projects can use
- Specifications:
  - Traffic policy
  - Traffic telemetry
  - Traffic Management
- Linkerd implements some of the above (e.g. TrafficSplit)



# Architecture



# "Meshing" pods

- To use Linkerd's features, pods need to have the linkerd-proxy sidecar
- Pods with a linkerd-proxy are "meshed" pods
- Add **linkerd.io/inject: enabled** annotation
- Most functionality at the client so ensure clients are "meshed"

```
apiVersion: v1
kind: Pod
metadata:
  name: debug
  labels:
    name: debug
  annotations:
    linkerd.io/inject: enabled
```

# Linkerd init container

- Linkerd adds an init container to your pod that modifies **iptables** rules
- This requires **NET\_ADMIN** capabilities
  - ➔ some clusters prohibit this
- Solution: use the Linkerd CNI plugin
  - ➔ <https://linkerd.io/2.11/features/cni>

# Load Balancing

- Kubernetes Load Balancing
  - Connection level
  - Issue with gRPC which uses HTTP/2; HTTP/2 does request multiplexing over a single connection
- Linkerd Load Balancing
  - Request level based on latency (EWMA)
  - Works well with gRPC out of the box
  - Client-side so client needs to be meshed

# Protocol Detection

- Built-in protocol detection
- HTTP/1.1, HTTP/2 or gRPC
  - For these, Linkerd provides full functionality
- What if it is another protocol?
  - Simply proxy the request as raw TCP
  - TLS is also processed as raw TCP
- You can mark ports as **opaque** to tell Linkerd to skip protocol detection  
config.linkerd.io/opaque-ports (see [TCP Proxying and Protocol Detection | Linkerd](#))



# Service Profiles

- Out of the box, **golden metrics** are at the service level
- You can obtain the metrics per **endpoint** of your service
  - via a **Service Profile**
- Service profile is extra config: more YAML!
  - or use CLI: can use OpenAPI/Swagger or protobuf file to learn about endpoints
  - or monitor traffic: **linkerd viz profile**
- They are not only for per route metrics
  - Retries and timeouts
- Optional ⚠

# Exercise

---

- Install Linkerd on AKS
- Install an application with meshed pods



# Ingress and Service Mesh

- Linkerd does not come with its own Ingress Controller
- Add your Ingress Controller of choice
  - NGINX, Traefik, ...
- Mesh the Ingress Controller to secure traffic to backend services and to obtain metrics
  - not possible for Ingress solutions with external proxies (e.g., App Gateway)

⚠ Beware of **sticky sessions**

# Exercise

---

- Install meshed Ingress Controller
- Configure the ingress resource



# Key Takeaways

---

## Cluster deployment

- Azure CNI
- Azure Firewall for egress control
- Network Policies for east-west traffic control

## Service Mesh

- Golden metrics
- Encryption of traffic in the cluster
- Improved load balancing
- Application-level features: retries, timeouts