



# Advanced pod concepts

---

Kubernetes Deep Dive

```
mirror_mod = modifier_ob.  
Set mirror object to mirror.  
mirror_mod.mirror_object =  
operation == "MIRROR_X":  
mirror_mod.use_x = True  
mirror_mod.use_y = False  
mirror_mod.use_z = False  
operation == "MIRROR_Y":  
mirror_mod.use_x = False  
mirror_mod.use_y = True  
mirror_mod.use_z = False  
operation == "MIRROR_Z":  
mirror_mod.use_x = False  
mirror_mod.use_y = False  
mirror_mod.use_z = True  
  
selection at the end -add  
ob.select= 1  
modifier_ob.select=1  
context.scene.objects.active  
("Selected" + str(modifier_ob.  
mirror_ob.select = 0  
= bpy.context.selected_object  
data.objects[one.name].select  
  
print("please select exactly  
  
-- OPERATOR CLASSES --  
  
types.Operator):  
X mirror to the selected  
object.mirror_mirror_x"  
mirror X"  
  
context):  
context.active_object is not
```

# What's in this module?

---

Pod lifecycle

Init  
Containers

Multi-  
container  
pods

Scheduling



# Pod Lifecycle

---



# Pod phases

## Pending

- accepted by cluster

## Running

- running on a node

## Succeeded

- terminated with success

## Failed

- terminated with failure

## Unknown

- indicative of node communication failure

# Container states

## Waiting

- still doing things to complete start up (e.g., pull)

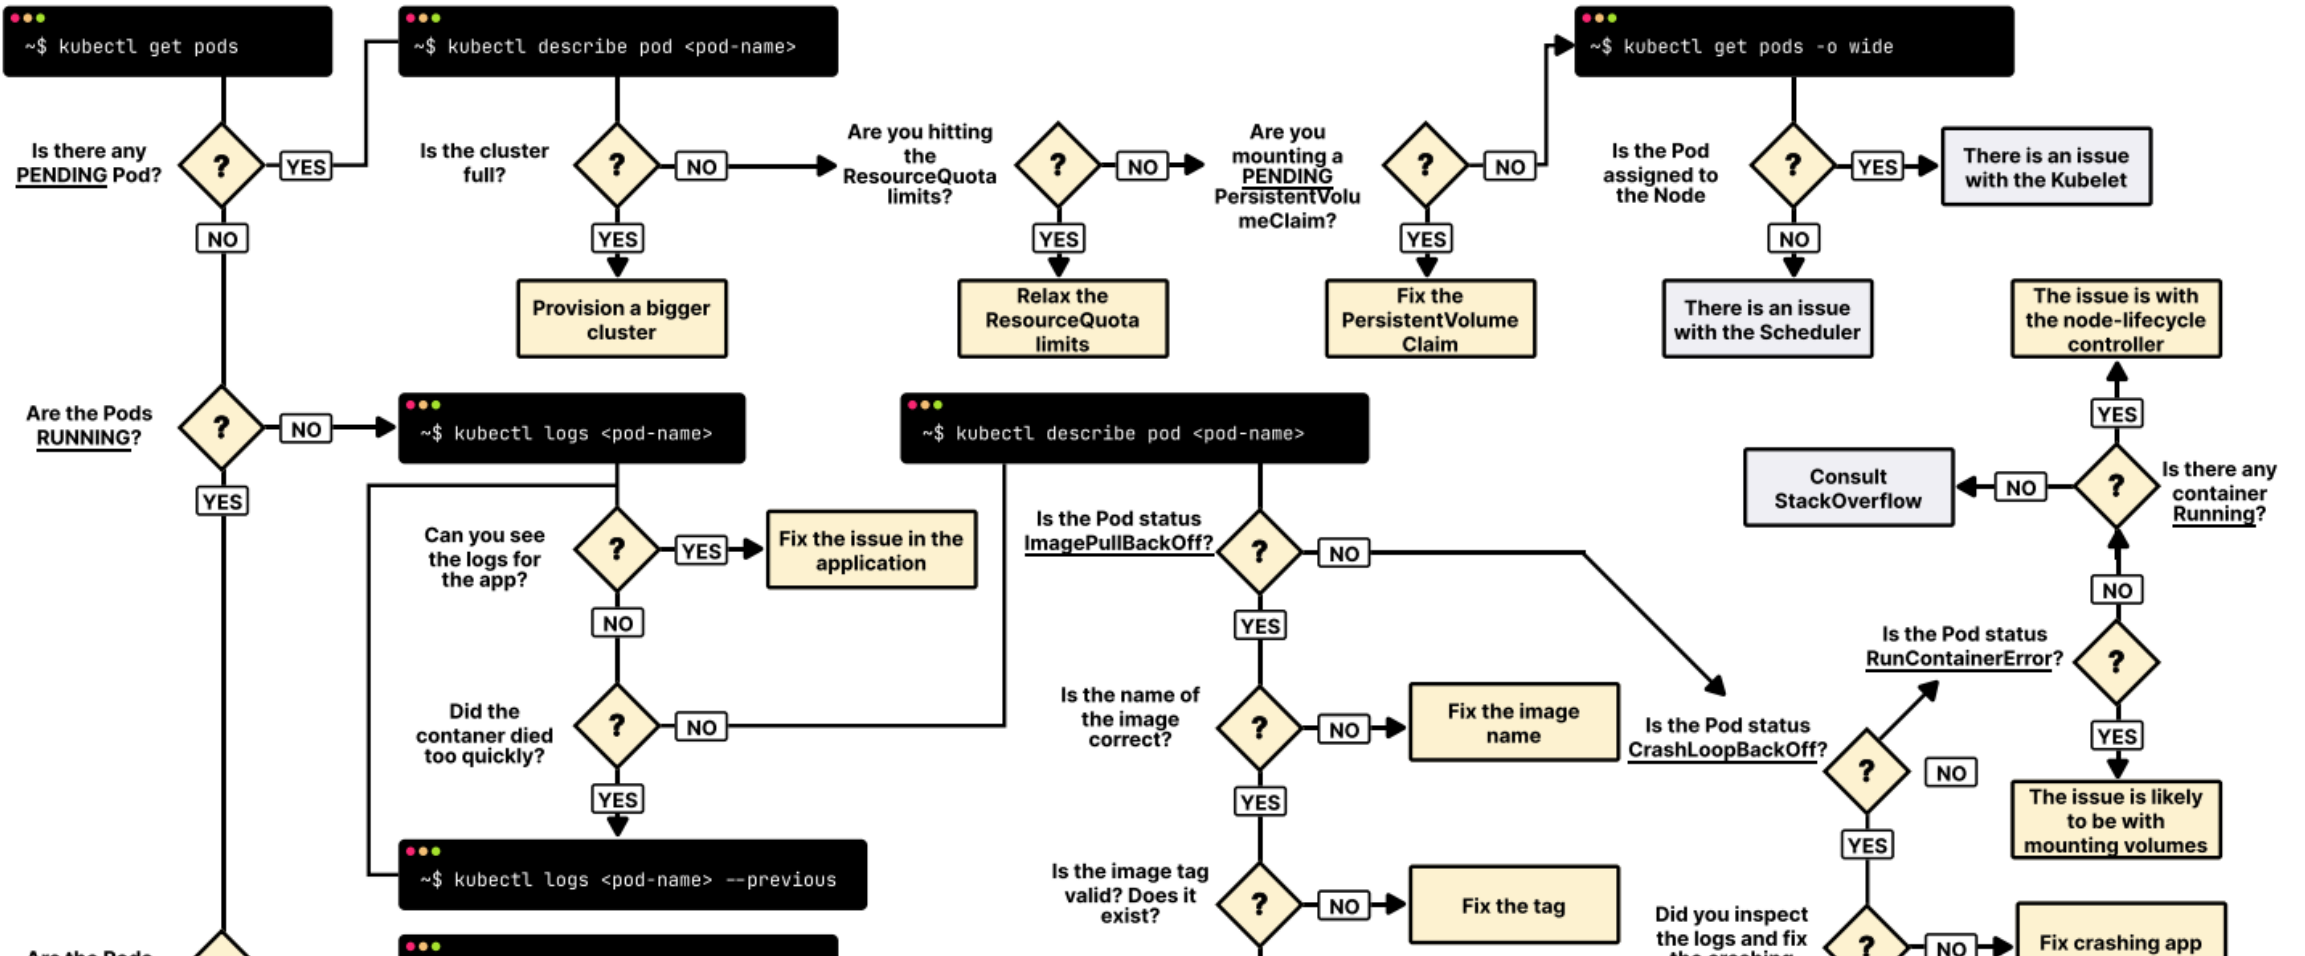
## Running

- container is running

## Terminated

- ran to completion or failed

# START





# Init containers

---



# Init Containers


---



## **Init containers run before the application containers**

run to completion

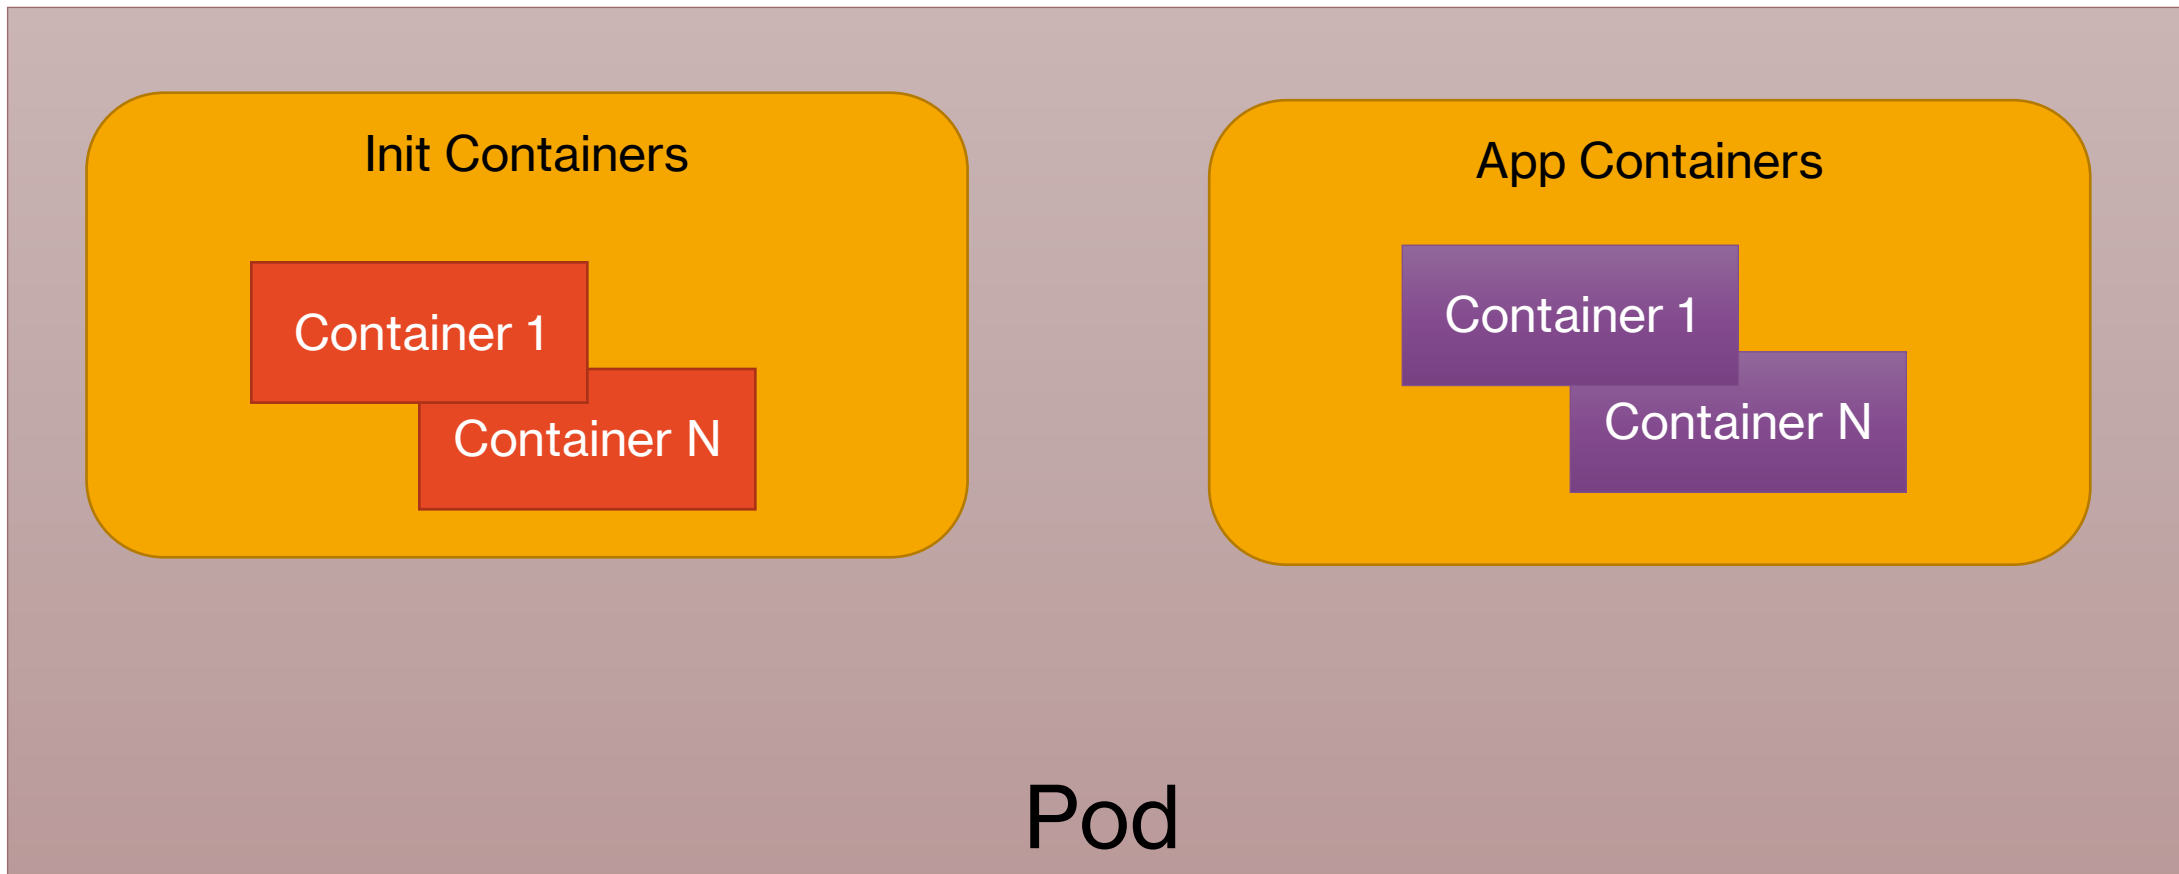
run one after the other



## **Each init container can have a separate image**

run any tool you want without needing to add it to app image





Init containers run sequentially  
**before** the app containers run



# Use cases

- Modify network rules
- Wait for events or sleep
- Register pods with remote systems
- Download files to a volume  
e.g. git clone



```
apiVersion: v1
kind: Pod
metadata:
  name: simple-init
  labels:
    name: simple-init
spec:
  restartPolicy: Never
  initContainers:
    - name: init-container1
      image: busybox
      command:
        - sh
        - -c
        - echo "Hello from init container 1"
    - name: init-container2
      image: busybox
      command:
        - sh
        - -c
        - exit 1
  containers:
    - name: simple-init
      image: nginx
      resources:
        limits:
          memory: "64Mi"
          cpu: "100m"
      ports:
        - containerPort: 80
```

# Example

- Two init containers that run sequentially
- Second init container has non-zero exit code
- kubelet will restart the second init container unless restartPolicy=Never
- **Important:** set requests/limits in case namespace has ResourceQuota

# Multi container pods

---





# Multi-container pods

- Run multiple co-located containers in a pod
- Containers run in parallel and can access shared resources
  - Storage: shared volumes
  - Network: shared network namespace (IP, network ports, ...)

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-container
  labels:
    name: multi-container
spec:
  containers:
  - name: container1
    image: k8s.gcr.io/pause:3.1
  - name: container2
    image: k8s.gcr.io/pause:3.1
  - name: container3
    image: k8s.gcr.io/pause:3.1
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-multi-deployment
  labels:
    app: nginx-multi
spec:
  replicas: 1
  selector:
    matchLabels:
      app: web-multi
  template:
    metadata:
      labels:
        app: web-multi
    spec:
      volumes:
        - name: webapp
          emptyDir:
            medium: Memory
      initContainers:
        - name: configure
          image: alpine/git
          volumeMounts:
            - name: webapp
              mountPath: /work
          command:
            - git
            - clone
            - https://github.com/gbaeke/static-web.git
            - '/work'
      containers:
        - name: pull
          image: alpine/git
          volumeMounts:
            - name: webapp
              mountPath: /work
          command:
            - "/bin/sh"
            - "-c"
            - "cd /work; while true; do git pull; sleep 5; done;"
          resources:
            limits:
              memory: 64Mi
              cpu: 100m
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
          volumeMounts:
```

# Example

- Init container to clone the repo
- Container to run **git pull**
- Main container **nginx** to serve the contents of the git repo



# Sidecars

---





# Sidecar

- One or more containers that run along the main container
- Support the main container without changing it
- Important:
  - configure health checks
  - set appropriate requests/limits

# Sidecar request and limits

```
spec:
  replicas: 1
  selector:
    matchLabels:
      app: web-multi
  template:
    metadata:
      labels:
        app: web-multi
    annotations:
      linkerd.io/inject: "enabled"
      config.linkerd.io/proxy-cpu-limit: "0.5"
      config.linkerd.io/proxy-memory-limit: 128Mi
```

- Especially important in combination with ResourceQuota
- Sidecar injectors might need additional configuration



# Ephemeral containers

---



# Ephemeral Containers

- Alpha feature in Kubernetes 1.22
- Run a container in a pod temporarily
- Useful for interactive troubleshooting
  - Container crashed
  - Containers without debugging tools or shells (scratch & distroless)
- How?
  - **kubectl debug -it name --image busybox --target targetpod**





# Disruptions

---





# Disruptions

- Voluntary versus involuntary
- Use **PodDisruptionBudget** to limit the number of pods down from involuntary disruptions
  - set the number of replicas you tolerate having, versus the intended number
  - used by tools that use the eviction API (e.g. kubectl drain)
- Example:
  - Replicas in spec: 10
  - PodDisruptionBudget: 6
  - Eviction API will allow disruption of 4 pods at a time

## Example

- Pods with label **web-multi** should have a minimum of 2 pods available
  - percentages are allowed
- If initial deployment specified 3, then 1 disruption is allowed

```
gbaeke ~ static-web master $ k get pdb
```

NAME	MIN AVAILABLE	MAX UNAVAILABLE	ALLOWED DISRUPTIONS	AGE
web-multi-pdb	2	N/A	1	15m



```
# use policy/v1 in k8s 1.21+
apiVersion: policy/v1beta1
kind: PodDisruptionBudget
metadata:
  name: web-multi-pdb
spec:
  minAvailable: 2
```

```
app: web-multi
```

# Topology spread constraints


---



# Pod spreading

- Control how pods are spread across a target **topology** of the cluster
  - Regions, zones, nodes, ...
- AKS uses the **topology.kubernetes.io/zone** label to:
  - Identify the zone of the node (if Availability Zones are used)
  - Distribute the pods across the zones

```
storagetier=Premium_LRS  
topology.kubernetes.io/region=westeurope  
topology.kubernetes.io/zone=westeurope-1
```



# Topology Spread Constraints

- New field in **pod spec**: topologySpreadConstraints
  - Applies to pods based on selector
- Topology is defined by a key of a node label
  - `kubectl get nodes --show-labels`
  - `kubectl describe nodes`
- Topology constraints at cluster level are **not possible** in a managed offering like AKS



```
spec:
  topologySpreadConstraints:
    - maxSkew: 1
      topologyKey: kubernetes.azure.com/agentpool
      whenUnsatisfiable: DoNotSchedule
      labelSelector:
        matchLabels:
          app: topo-demo
```

# maxSkew


Maximum permitted difference between # of matching pods in the target topology and global minimum



skew would be 2




maxSkew = 1





# maxSkew

Maximum permitted difference between # of matching pods in the target topology and global minimum





**Now onto...**

---

dapn